

Learning Linux Device Drivers

Jeremiah Mahler (jmmahler@gmail.com)

August 10, 2013

Contents

1	Introduction	2
2	Hello, World	2
2.1	hello	2
2.2	param	4
3	Read/Write Data	5
3.1	data_chr	5
3.2	data_rw	7
3.3	data_sk	7
3.4	ioctlx	7
3.5	null	7
3.6	zero	7
4	Sysfs	7
4.1	sysx_file	7
4.2	sysx_file2	7
4.3	sysx_group	7
4.4	sysx_ktype	7
4.5	sysx_ktype2	7
5	Concurrency	7
5.1	fifo_rw	7
5.2	fifo_sysfs	7
5.3	fifo_xxx	7
5.4	fifo_fix	7

1 Introduction

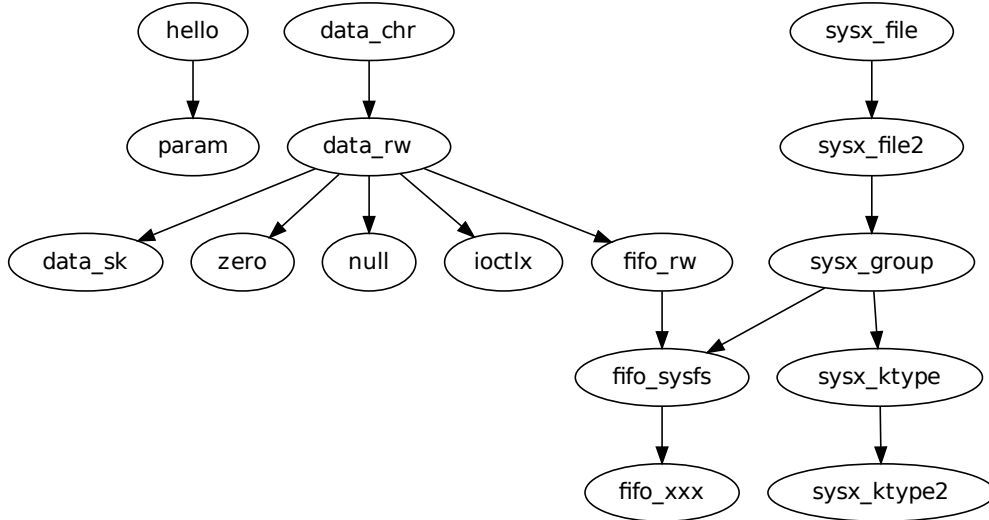


Figure 1: Hierarchy of kernel module examples. Simplest at the top downward to the more complex.

There are many excellent books about Linux device drivers^{123.4} However, in this authors experience, they were difficult to learn from. It certainly was not from their lack of detail. Clearly each of the authors have a profound understanding of the Linux kernel and their books reflect this. If anything it is due to this lack of simplicity.

The drivers described in this document aim to be simple and concise. Each one introduces as few concepts as possible. And each driver is a fully working example⁵. Many of the drivers are built in stages. Each stage introduces a new concept. And the changes can be concisely described showing the differences (`diff`). Figure 1 shows the hierarchy of driver examples.

2 Hello, World

2.1 hello

The `hello` module (Listing 1) simply prints message when it is loaded and unload.

```
hello$ make
(should compile without error, resulting in hello.ko)
hello$ sudo insmod hello.ko
Hello, World
hello$ sudo rmmod hello
Goodbye, cruel world
```

¹J. Corbet, A. Rubini, and G. Kroah-Hartman. *Linux Device Drivers*. O'Reilly Media, 2009. ISBN: 9780596555382.

²S. Venkateswaran. *Essential Linux Device Drivers*. Pearson Education, 2008. ISBN: 9780132715812.

³R. Love. *Linux Kernel Development*. Developer's Library. Pearson Education, 2010. ISBN: 9780768696790.

⁴Robert. Love. *Linux System Programming: Talking Directly to the Kernel and C Library*. O'Reilly Media, 2013. ISBN: 9781449341541.

⁵If you think a driver might behave in a different way if you changed something, try it and see. This is a far better way to solidify your understanding than just following the authors examples.

```

1  #include <linux/init.h>
2  #include <linux/module.h>
3
4  static int __init hello_init(void)
5  {
6      printk(KERN_ALERT "Hello , World\n" );
7      return 0;
8  }
9
10 static void __exit hello_exit(void)
11 {
12     printk(KERN_ALERT "Goodbye, cruel world\n" );
13 }
14
15 MODULE_AUTHOR(" Jeremiah Mahler <jmmahler@gmail.com>" );
16 MODULE_LICENSE("GPL" );
17
18 module_init( hello_init );
19 module_exit( hello_exit );

```

Listing 1: Hello, World module in hello/hello.c

The `module_init` (line 18) and `module_exit` (line 19) tell the kernel which functions to call when this module is loaded (`insmod`) and unloaded (`rmmod`).

The `__init` (line 4) and `__exit` (line 10) are optional hints for the compiler. For example in the case of `__init`, this tells the kernel that it may discard the code after initialization has been completed.

Both the init function (line 4) and the exit function (line 10) are declared **static**. Since these functions are not meant to be used outside the scope of this file, declaring them **static** enforces this constraint.⁶

The `printk` statements are the `printf` of the kernel domain. There are various levels, in this case `KERN_ALERT` is used which will cause the messages to appear on the console. Notice that there is no comma between the level and the message.

The `MODULE_AUTHOR` and `MODULE_LICENSE` on lines 15 and 16 are optional but recommended.⁷ There are various other `MODULE_*` options as well (`linux/module.h`).

⁶Corbet, Rubini, and Kroah-Hartman, see n. 1, Pg. 52.

⁷Ibid., Pg. 51.

2.2 param

The `param` module expands upon the `hello` module to take a parameter specifying how many times to print the message.

```
param$ sudo insmod hello.ko howmany=2
Hello, World
Hello, World
param$ sudo rmmod hello
Goodbye, cruel world
Goodbye, cruel world
```

Listing 2 shows the differences between this parameterized hello world module and the previous `hello` module.

```
1  --- ../hello/hello.c      2013-08-09 12:23:58.222416131 -0700
2  +++ hello.c 2013-08-09 12:53:38.082434726 -0700
3  @@ -1,15 +1,28 @@
4  #include <linux/init.h>
5  #include <linux/module.h>
6  +#include <linux/moduleparam.h>
7  +
8  +static int howmany = 1;
9  +module_param(howmany, int, S_IRUGO);
10
11 static int __init hello_init(void)
12 {
13 - printk(KERN_ALERT "Hello , World\n");
14 + int i;
15 +
16 + for (i = 0; i < howmany; i++) {
17 +     printk(KERN_ALERT "Hello , World\n");
18 + }
19 +
20     return 0;
21 }
22
23 static void __exit hello_exit(void)
24 {
25 - printk(KERN_ALERT "Goodbye, cruel world\n");
26 + int i;
27 +
28 + for (i = 0; i < howmany; i++) {
29 +     printk(KERN_ALERT "Goodbye, cruel world\n");
30 + }
31 }
32
33 MODULE_AUTHOR("Jeremiah Mahler <jmmahler@gmail.com>");
```

Listing 2: `param$ diff -u hello.c ../hello/hello.c`

To use a parameter a global variable has been created named `howmany` on line 8. And on line 9 the `module_param` function is used to tell the kernel about this parameter ⁸.

On lines 13-19 and 25-30 it can be seen that the same message is printed `howmany` times.

⁸The `module_param` function create a `sysfs` entry in `/sys/module/parameters/howmany`. `sysfs` will be discussed in detail in later modules.

3 Read/Write Data

The `data` module allocates a memory from ram which can be read from and written to. This is accomplished as a character device and supports all the usual file operations.

3.1 `data_chr`

The first step is to construct the basic infrastructure for a character driver as shown in Listing 3. It doesn't do anything useful but it will simplify the description of upcoming drivers.

The `DEVICE_NAME` (line 8) is just a shortcut for the name which is used in several places.

Lines 10-17 are the global variables that will be used. The `struct data_dev` is the per device structure. Notice that a character device is placed inside.

The `file_operations` (line 19-21) in this case only defines the `.owner`. Upcoming modules will add references to the `open`, `close`, `read`, `write`, and `seek` functions to this structure.

`alloc_chrdev_region` (line 26) allocates a major and minor number for the character device.⁹ In this case only one major and minor pair is needed.

Functions such as `alloc_chrdev_region` may fail and when they do anything that has been created up to that point must be undone to ensure the kernel is left in a consistent state. A common way this is done is using `goto` statements which branch to different steps in the exit sequence.¹⁰ It can be seen that if `alloc_chrdev_region` fails its `goto` (line 29) will branch to line 66. Since nothing was created up to that point nothing has to be undone.

`class_create` (line 32) establishes a "class" for this module which is also represented in sysfs under `/sys/class/data`. This object will be used later as an argument to `device_create`.

Since the per device structure is just a pointer it must be allocated before it is used (line 32).

To establish the character device it must be initialized and added (lines 41-43). And finally the device is created (line 49). This device will now appear under `/dev/fifo0`.

```
1  #include <linux/cdev.h>
2  #include <linux/device.h>
3  #include <linux/fs.h>
4  #include <linux/module.h>
5  #include <linux/slab.h>
6  #include <linux/uaccess.h>
7
8  #define DEVICENAME "data"
9
10 static dev_t data_major;
11 struct class *data_class;
12 struct device *data_device;
13
14 struct data_dev {
15     struct cdev cdev;
16 } *data_devp;
17
18 struct file_operations data_fops = {
19     .owner = THIS_MODULE,
20 };
21
```

⁹Corbet, Rubini, and Kroah-Hartman, see n. 1, Pg. 66.

¹⁰Ibid., Pg. 53.

```

22 static int __init data_init(void)
23 {
24     int err = 0;
25
26     err = alloc_chrdev_region(&data_major, 0, 1, DEVICENAME);
27     if (err < 0) {
28         printk(KERN_WARNING "Unable to register device\n");
29         goto err_chrdev_region;
30     }
31
32     data_class = class_create(THIS_MODULE, DEVICENAME);
33
34     data_devp = kmalloc(sizeof(struct data_dev), GFP_KERNEL);
35     if (!data_devp) {
36         printk(KERN_WARNING "Unable to kmalloc data_devp\n");
37         err = -ENOMEM;
38         goto err_malloc_data_devp;
39     }
40
41     cdev_init(&data_devp->cdev, &data_fops);
42     data_devp->cdev.owner = THIS_MODULE;
43     err = cdev_add(&data_devp->cdev, data_major, 1);
44     if (err) {
45         printk(KERN_WARNING "cdev_add failed\n");
46         goto err_cdev_add;
47     }
48
49     data_device = device_create(data_class, NULL,
50                               MKDEV(MAJOR(data_major), 0), NULL, "data%d", 0);
51     if (IS_ERR(data_device)) {
52         printk(KERN_WARNING "device_create failed\n");
53         err = PTR_ERR(data_device);
54         goto err_device_create;
55     }
56
57     return 0; /* success */
58
59 err_device_create:
60     cdev_del(&data_devp->cdev);
61 err_cdev_add:
62     kfree(data_devp);
63 err_malloc_data_devp:
64     class_destroy(data_class);
65     unregister_chrdev_region(data_major, 1);
66 err_chrdev_region:
67
68     return err;
69 }
70
71 static void __exit data_exit(void)
72 {
73     device_destroy(data_class, data_major);
74
75     cdev_del(&data_devp->cdev);
76

```

```

77     kfree (data_devp);
78
79     class_destroy (data_class);
80
81     unregister_chrdev_region (data_major , 1);
82 }
83
84 MODULE_AUTHOR("Jeremiah Mahler <jmmahler@gmail.com>");
85 MODULE_LICENSE("GPL");
86
87 module_init (data_init);
88 module_exit (data_exit);

```

Listing 3: data_chr driver.

3.2 data_rw

3.3 data_sk

3.4 ioctlx

3.5 null

3.6 zero

4 Sysfs

4.1 sysx_file

4.2 sysx_file2

4.3 sysx_group

4.4 sysx_ktype

4.5 sysx_ktype2

5 Concurrency

5.1 fifo_rw

5.2 fifo_sysfs

5.3 fifo_xxx

5.4 fifo_fix

References

Corbet, J., A. Rubini, and G. Kroah-Hartman. *Linux Device Drivers*. O'Reilly Media, 2009. ISBN: 9780596555382.

Love, R. *Linux Kernel Development*. Developer's Library. Pearson Education, 2010. ISBN: 9780768696790.

Love, Robert. *Linux System Programming: Talking Directly to the Kernel and C Library*. O'Reilly Media, 2013. ISBN: 9781449341541.

Venkateswaran, S. *Essential Linux Device Drivers*. Pearson Education, 2008. ISBN: 9780132715812.