

Linux Module Programming in Small Diffs

Jeremiah Mahler (jimmahler@gmail.com)

August 9, 2013

Contents

1	Introduction	2
2	Hello, World	2
2.1	hello	2
2.2	param	4
3	Read/Write Data	5
3.1	data_chr	5
3.2	data_rw	7
3.3	data_sk	7
3.4	ioctlx	7
3.5	null	7
3.6	zero	7
4	Sysfs	7
4.1	sysx_file	7
4.2	sysx_file2	7
4.3	sysx_group	7
4.4	sysx_ktype	7
4.5	sysx_ktype2	7
5	Concurrency	7
5.1	fifo_rw	7
5.2	fifo_sysfs	7
5.3	fifo_xxx	7
5.4	fifo_fix	7

1 Introduction

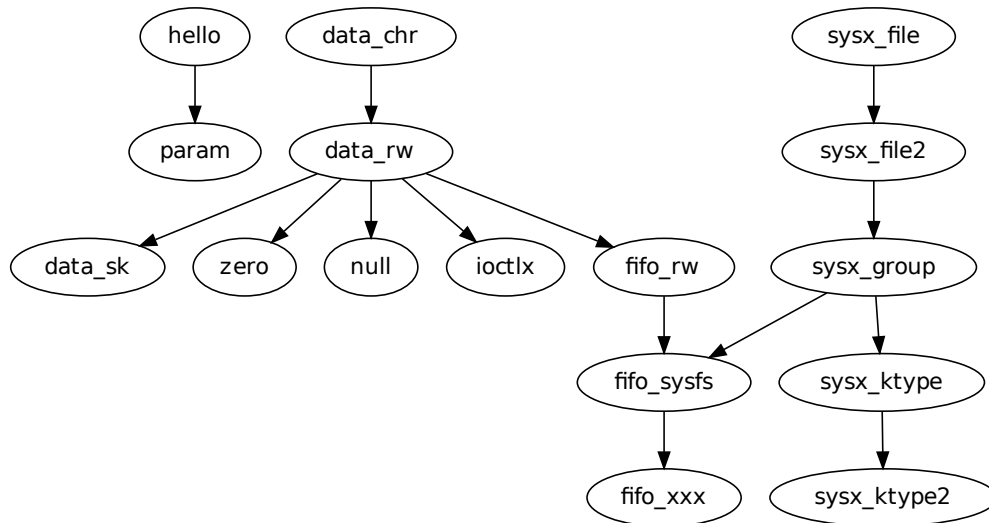


Figure 1: Hierarch of kernel module examples.

2 Hello, World

2.1 hello

The `hello` module (Listing 1) simply prints message when it is loaded and unload.

```
hello$ make
(should compile without error, resulting in hello.ko)
hello$ sudo insmod hello.ko
Hello, World
hello$ sudo rmmod hello
Goodbye, cruel world
```

The `__init` and `__exit` on lines 4 and 10 are optional hints for the compiler. For example in the case of `__init` it may discard that code after initialization has been completed.

The `printk` statements are the `printf` of the kernel domain. There are various levels, in this case `KERN_ALERT` is used which will cause the messages to appear on the console. Notice that there is no comma between the level and the message.

The `MODULE_AUTHOR` and `MODULE_LICENSE` on lines 15 and 16 are optional but recommended. There are various other `MODULE_*` as well (`linux/module.h`).

The `module_init` and `module_exit` tell the kernel which functions to call when this module is loaded (`insmod`) and unloaded (`rmmod`).

```

1 #include <linux/init.h>
2 #include <linux/module.h>
3
4 static int __init hello_init(void)
5 {
6     printk(KERN_ALERT "Hello , World\n");
7     return 0;
8 }
9
10 static void __exit hello_exit(void)
11 {
12     printk(KERN_ALERT "Goodbye, cruel world\n");
13 }
14
15 MODULE_AUTHOR("Jeremiah Mahler <jmmahler@gmail.com>");
16 MODULE_LICENSE("GPL");
17
18 module_init( hello_init );
19 module_exit( hello_exit );

```

Listing 1: Hello, World module in hello/hello.c

2.2 param

The `param` module expands upon the `hello` module to take a parameter specifying how many times to print the message.

```
param$ sudo insmod hello.ko howmany=2
Hello, World
Hello, World
param$ sudo rmmod hello
Goodbye, cruel world
Goodbye, cruel world
```

Listing 2 shows the differences between this parameterized hello world module and the previous `hello` module.

```
1  --- ../hello/hello.c      2013-08-09 12:23:58.222416131 -0700
2  +++ hello.c 2013-08-09 12:53:38.082434726 -0700
3  @@ -1,15 +1,28 @@
4  #include <linux/init.h>
5  #include <linux/module.h>
6  +#include <linux/moduleparam.h>
7  +
8  +static int howmany = 1;
9  +module_param(howmany, int, SIRUGO);
10
11 static int __init hello_init(void)
12 {
13 - printk(KERN_ALERT "Hello , World\n");
14 + int i;
15 +
16 + for (i = 0; i < howmany; i++) {
17 +     printk(KERN_ALERT "Hello , World\n");
18 + }
19 +
20     return 0;
21 }
22
23 static void __exit hello_exit(void)
24 {
25 - printk(KERN_ALERT "Goodbye, cruel world\n");
26 + int i;
27 +
28 + for (i = 0; i < howmany; i++) {
29 +     printk(KERN_ALERT "Goodbye, cruel world\n");
30 + }
31 }
32
33 MODULE_AUTHOR(" Jeremiah Mahler <jmmahler@gmail.com>");
```

Listing 2: `param$ diff -u hello.c ../hello/hello.c`

To use a parameter a global variable has been created named `howmany` on line 8. And on line 9 the `module_param` function is used to tell the kernel about this parameter ¹.

On lines 13-19 and 25-30 it can be seen that the same message is printed `howmany` times.

¹The `module_param` function create a `sysfs` entry in `/sys/module/parameters/howmany`. `sysfs` will be discussed in detail in later modules.

3 Read/Write Data

The `data` module allocates a memory from ram which can be read from and written to. This is accomplished as a character device and supports all the usual file operations.

3.1 `data_chr`

The first step is to construct the basic infrastructure for a character driver as shown in Listing 3.

The `DEVICE_NAME` on line 8 defines the string which will be used to define the module name in later functions.

Lines 10-17 are the global variables that will be used. The `struct data_dev` is the per device structure. Notice that a character device is placed inside.

The `file_operations` (line 19-21) in this case only defines the `.owner`. Upcoming modules will and references to the `open`, `close`, `read`, `write`, and `seek` functions to this structure.

The `data_cleanup` function takes care of unregistering and removing the various that were created during `data_init`. It will be called if module initialization fails or during module removal. Different authors use different styles for cleanup. It is not always done with a separate function. Some use several `goto` labels each with levels of items to remove ².

```
1  #include <linux/cdev.h>
2  #include <linux/device.h>
3  #include <linux/fs.h>
4  #include <linux/module.h>
5  #include <linux/slab.h>
6  #include <linux/uaccess.h>
7
8  #define DEVICENAME "data"
9
10 static dev_t data_major;
11 static int cdev_add_done;
12 struct class *data_class;
13 struct device *data_device;
14
15 struct data_dev {
16     struct cdev cdev;
17 } *data_devp;
18
19 struct file_operations data_fops = {
20     .owner = THIS_MODULE,
21 };
22
23 static void data_cleanup(void)
24 {
25     if (data_major) {
26         unregister_chrdev_region(data_major, 1);
27     }
28
29     if (data_device) {
30         device_destroy(data_class, data_major);
31     }
32
33     if (cdev_add_done) {
34         cdev_del(&data_devp->cdev);
35     }
36
```

²Recalling the `__init` directive, it might be beneficial use the `goto` style since this code will be purged after init.

```

37     if (data_devp) {
38         kfree(data_devp);
39     }
40
41     if (data_class) {
42         class_destroy(data_class);
43     }
44 }
45
46 static int __init data_init(void)
47 {
48     int err = 0;
49
50     data_major = 0;
51     data_class = NULL;
52     data_device = NULL;
53     data_devp = NULL;
54     cdev_add_done = 0;
55
56     if (alloc_chrdev_region(&data_major, 0, 1, DEVICENAME) < 0) {
57         printk(KERN_WARNING "Unable to register device\n");
58         err = -1;
59         goto err_out;
60     }
61
62     /* populate sysfs entries */
63     /* /sys/class/data/data0/ */
64     data_class = class_create(THIS_MODULE, DEVICENAME);
65
66     data_devp = kmalloc(sizeof(struct data_dev), GFP_KERNEL);
67     if (!data_devp) {
68         printk(KERN_WARNING "Unable to kmalloc data_devp\n");
69         err = -ENOMEM;
70         goto err_out;
71     }
72
73     cdev_init(&data_devp->cdev, &data_fops);
74     data_devp->cdev.owner = THIS_MODULE;
75     err = cdev_add(&data_devp->cdev, data_major, 1);
76     if (err) {
77         printk(KERN_WARNING "cdev_add failed\n");
78         goto err_out;
79     } else {
80         cdev_add_done = 1;
81     }
82
83     /* send uevents to udev, so it'll create /dev nodes */
84     /* /dev/data0 */
85     data_device = device_create(data_class, NULL,
86                               MKDEV(MAJOR(data_major), 0), NULL, "data%d", 0);
87
88     return 0; /* success */
89
90 err_out:
91     data_cleanup();

```

```

92     return err;
93 }
94
95 static void __exit data_exit(void)
96 {
97     data_cleanup();
98 }
99
100 MODULE_AUTHOR("Jeremiah Mahler <jmmahler@gmail.com>");
101 MODULE_LICENSE("GPL");
102
103 module_init(data_init);
104 module_exit(data_exit);

```

Listing 3: Data driver infrastructure.

3.2 data_rw

3.3 data_sk

3.4 ioctlx

3.5 null

3.6 zero

4 Sysfs

4.1 sysx_file

4.2 sysx_file2

4.3 sysx_group

4.4 sysx_ktype

4.5 sysx_ktype2

5 Concurrency

5.1 fifo_rw

5.2 fifo_sysfs

5.3 fifo_xxx

5.4 fifo_fix

References

- Corbet, J., A. Rubini, and G. Kroah-Hartman. *Linux Device Drivers*. O'Reilly Media, 2009. ISBN: 9780596555382.
- Love, R. *Linux Kernel Development*. Developer's Library. Pearson Education, 2010. ISBN: 9780768696790.
- Love, Robert. *Linux System Programming: Talking Directly to the Kernel and C Library*. O'Reilly Media, 2013. ISBN: 9781449341541.
- Venkateswaran, S. *Essential Linux Device Drivers*. Pearson Education, 2008. ISBN: 9780132715812.